

Importation des Libraires

```
In [80]: import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
from wordcloud import WordCloud, STOPWORDS
import re
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.linear_model import LogisticRegression
from sklearn.svm import LinearSVC
from sklearn.ensemble import RandomForestClassifier, GradientBoostingClassifier
from sklearn.naive_bayes import GaussianNB
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import accuracy_score, classification_report, confusion_matrix, f1_score, roc_auc_score, recall_score, precision_score
from sklearn.model_selection import train_test_split
import xgboost as xgb
import lightgbm as lgb
from sentence_transformers import SentenceTransformer
from sklearn.model_selection import cross_val_score
import joblib
```

1-Exploration et Analyse des Données (EAD)

1. Analyse des colonnes et distribution des classes

```
In [81]: train_data = pd.read_csv('train.En.csv')
train_data.info()
train_data.head()
train_data.describe()
counts = train_data['sarcastic'].value_counts()


<class 'pandas.core.frame.DataFrame'>
RangeIndex: 3468 entries, 0 to 3467
Data columns (total 10 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   Unnamed: 0        3468 non-null    int64  
 1   tweet             3467 non-null    object  
 2   sarcastic         3468 non-null    int64  
 3   rephrase          867 non-null    object  
 4   sarcasm           867 non-null    float64 
 5   irony              867 non-null    float64 
 6   satire             867 non-null    float64 
 7   understatement     867 non-null    float64 
 8   overstatement      867 non-null    float64 
 9   rhetorical_question 867 non-null    float64 
dtypes: float64(6), int64(2), object(2)
memory usage: 271.1+ KB
```

1. Mesure de la longueur des textes


```
In [82]: train_data['word_count'] = train_data['tweet'].apply(lambda x: len(str(x).split()))
train_data['char_count'] = train_data['tweet'].apply(lambda x: len(str(x)))
```

1. Visualisation des données


```
In [83]: plt.hist(train_data['word_count'], bins=30)
plt.title('Distribution des longueurs de mots')
plt.xlabel('Nombre de mots')
plt.ylabel('Fréquence')
plt.show()
```




```
In [84]: plt.hist(train_data['char_count'], bins=30)
plt.title('Distribution des longueurs de caractères')
plt.xlabel('Nombre de caractères')
plt.ylabel('Fréquence')
plt.show()
```



```
In [85]: plt.figure(figsize=(6,4))
counts.plot(kind='bar', color=['skyblue', 'salmon'])
plt.xticks([0,1], ['Non sarcastique', 'Sarcastique'], rotation=0)
plt.title("Répartition des tweets sarcastiques")
plt.ylabel("Nombre de tweets")
for i, v in enumerate(counts.values):
    plt.text(i, v + 5, str(v), ha='center')
plt.show()
```



```
In [86]: tweets_sarcastic = ' '.join(train_data[train_data['sarcastic'] == 1]['tweet'].astype(str))
tweets_non_sarcastic = ' '.join(train_data[train_data['sarcastic'] == 0]['tweet'].astype(str))
stopwords=set(STOPWORDS)
wordcloud_sarcastic = WordCloud(width=800, height=400, background_color='white', stopwords=stopwords).generate(tweets_sarcastic)
wordcloud_non_sarcastic = WordCloud(width=800, height=400, background_color='white', stopwords=stopwords).generate(tweets_non_sarcastic)
plt.figure(figsize=(14, 6))
plt.subplot(1, 2, 1)
plt.imshow(wordcloud_sarcastic, interpolation='bilinear')
plt.axis('off')
plt.title('Word Cloud - Tweets Sarcastiques')
plt.subplot(1, 2, 2)
plt.imshow(wordcloud_non_sarcastic, interpolation='bilinear')
plt.axis('off')
plt.title('Word Cloud - Tweets Non Sarcastiques')
plt.show()
```



2-Prétraitement du texte

```
In [87]: train_data=train_data[["tweet","sarcastic"]]
train_data.head()
```

Out[87]:

		tweet	sarcastic
0	The only thing I got from college is a caffen...	1	
1	I love it when professors draw a big question ...	1	
2	Remember the hundred emails from companies whe...	1	
3	Today my pop-pop told me I was not “forced” to...	1	
4	@VolphanCarol @littlewhitty @mysticalmanatee I...	1	

1. Nettoyage et transformation du texte

```
In [88]: def clean_text(text):
    if not isinstance(text, str):
        return ""
    text = text.lower()
    text = re.sub(r"http\S+|www\S+", "", text)
    text = re.sub(r"@w+", "", text)
    text = re.sub(r"#", "", text)
    text = re.sub(r"[^a-zA-Z\s]", "", text)
    text = re.sub(r"\s+", " ", text).strip()
    words = [word for word in text.split() if word not in stopwords]
    text = " ".join(words)

    return text
```

```
In [89]: train_data['text_cleaned'] = train_data['tweet'].apply(clean_text)
train_data.head()
```

Out[89]:

	tweet	sarcastic	text_cleaned
0	The only thing I got from college is a caffeine addiction	1	thing got college caffeine addiction
1	I love it when professors draw a big question mark next to it	1	love professors draw big question mark next to it
2	Remember the hundred emails from companies who started covid	1	remember hundred emails companies covid starte...
3	Today my pop-pop told me I was not “forced” to go to college	1	today poppop told forced go college okay sure ...
4	@VolphanCarol @littlewhitty @mysticalmanatee I...	1	reported cancun cruz worrying heartbeats const...

1. Sélection des colonnes pertinentes

```
In [90]: train_data=train_data[['text_cleaned','sarcastic']]
train_data.head()
```

Out[90]:

	text_cleaned	sarcastic
0	thing got college caffeine addiction	1
1	love professors draw big question mark next to it	1
2	remember hundred emails companies covid starte...	1
3	today poppop told forced go college okay sure ...	1
4	reported cancun cruz worrying heartbeats const...	1

1. Séparation des features et du label

```
In [91]: X = train_data['text_cleaned'].astype(str)
y = train_data['sarcastic']
```

1. Division des données

```
In [92]: X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=42, stratify=y
)
```

3-Vectorisation (*transformation en nombres*)

1. Encodage des textes avec SBERT

```
In [93]: X_test_list = X_test.astype(str).tolist()
X_train_list = X_train.astype(str).tolist()

model_sbert = SentenceTransformer('all-MiniLM-L6-v2')
X_train_emb = model_sbert.encode(X_train_list, batch_size=32, show_progress_bar=True)
X_test_emb = model_sbert.encode(X_test_list, batch_size=32, show_progress_bar=True)

Batches: 100%|██████████| 87/87 [00:05<00:00, 16.09it/s]
Batches: 100%|██████████| 22/22 [00:01<00:00, 14.79it/s]
```

4-Application d'algorithmes de Machine Learning

1. Fonction d'entraînement ,d'évaluation des modèles

```
In [94]: def train_evaluate(model, X_train, y_train, X_test, y_test, name="Model"):
    print(f"\n===== {name} =====")
    model.fit(X_train, y_train)
    y_pred = model.predict(X_test)
    print(classification_report(y_test, y_pred, digits=3))

    cm = confusion_matrix(y_test, y_pred)
    sns.heatmap(cm, annot=True, fmt="d", cmap="Blues")
    plt.title(f"{name} - Confusion Matrix")
    plt.xlabel("Prédictions")
    plt.ylabel("Vérités terrain")
    plt.show()
    return y_pred
```

1. Entraînement et évaluation de tous les modèles


Il s'agit de:

- 1)Entraîner le modèle
- 2)Évaluer les performances : accuracy, precision, recall,F1-score
- 3)Afficher la matrice de confusion

(a) Modèle Logistic Regression

```
In [95]: logreg = LogisticRegression(max_iter=1000, class_weight='balanced')
f1_logreg = train_evaluate(logreg, X_train_emb, y_train, X_test_emb, y_test,
                            "Logistic Regression")
```

	precision	recall	f1-score	support
0	0.785	0.618	0.692	521
1	0.299	0.491	0.372	173
accuracy			0.586	694
macro avg	0.542	0.555	0.532	694
weighted avg	0.664	0.586	0.612	694




(b) Modèle SVM (Support Vector Classifier)

```
In [96]: svm = LinearSVC( class_weight='balanced', random_state=42)
f1_svm = train_evaluate(svm, X_train_emb, y_train, X_test_emb, y_test, "SVM")
```

===== SVM =====


	precision	recall	f1-score	support
0	0.783	0.618	0.691	521
1	0.297	0.486	0.368	173
accuracy			0.585	694
macro avg	0.540	0.552	0.530	694
weighted avg	0.662	0.585	0.611	694



(c) Modèle Random Forest

```
In [97]: rf = RandomForestClassifier( class_weight='balanced', random_state=42)
f1_rf = train_evaluate(rf, X_train_emb, y_train, X_test_emb, y_test, "Random Forest")
```

===== Random Forest =====				
	precision	recall	f1-score	support
0	0.750	0.998	0.857	521
1	0.000	0.000	0.000	173
accuracy			0.749	694
macro avg	0.375	0.499	0.428	694
weighted avg	0.563	0.749	0.643	694




(d) Modèle Naïve Bayes (GaussianNB)

```
In [98]: gnb = GaussianNB()
f1_gnb = train_evaluate(gnb, X_train_emb, y_train, X_test_emb, y_test, "Naive Bayes (GaussianNB)")
```

```
===== Naive Bayes (GaussianNB) =====
      precision    recall   f1-score   support
0         0.776     0.747     0.761     521
1         0.316     0.353     0.333     173

accuracy                           0.648     694
macro avg       0.546     0.550     0.547     694
weighted avg    0.662     0.648     0.655     694
```




(e) Modèle K-Nearest Neighbors (KNN)

```
In [99]: knn = KNeighborsClassifier(n_neighbors=5)
f1_knn = train_evaluate(knn, X_train_emb, y_train, X_test_emb, y_test, "KNN")
```

===== KNN =====

	precision	recall	f1-score	support
0	0.754	0.919	0.829	521
1	0.288	0.098	0.147	173
accuracy			0.715	694
macro avg	0.521	0.509	0.488	694
weighted avg	0.638	0.715	0.659	694



(f) Modèle XGBoost

```
In [100]: neg = sum([1 for label in y_train if label==0])
pos = sum([1 for label in y_train if label==1])
scale_pos_weight = neg / pos


xgb = xgb.XGBClassifier(
    n_estimators=200,
    learning_rate=0.1,
    max_depth=5,
    scale_pos_weight=scale_pos_weight,
    use_label_encoder=False,
    eval_metric='logloss',
    random_state=42
)
f1_xgb = train_evaluate(xgb, X_train_emb, y_train, X_test_emb, y_test, "XGBoost")
```

===== XGBoost =====

```
c:\Users\brahi\AppData\Local\Programs\Python\Python313\Lib\site-packages\xgboost\training.py:199: UserWarning: [22:14:38] WARNING: C:\actions-runner\_work\xgboost\xgboost\src\learner.cc:790:  
Parameters: { "use_label_encoder" } are not used.
```

```
bst.update(dtrain, iteration=i, fobj=obj)
```

	precision	recall	f1-score	support
0	0.757	0.914	0.828	521
1	0.308	0.116	0.168	173
accuracy			0.715	694
macro avg	0.532	0.515	0.498	694
weighted avg	0.645	0.715	0.663	694




(g) Modèle LightGBM

```
In [101]: import lightgbm as lgb

lgb_model = lgb.LGBMClassifier(
    objective='binary',
    class_weight='balanced',
    n_estimators=100,
    force_col_wise=True
)
f1_lgb = train_evaluate(lgb_model, X_train_emb, y_train, X_test_emb, y_test,
"LightGBM")
```

```
===== LightGBM =====
[LightGBM] [Info] Number of positive: 694, number of negative: 2080
[LightGBM] [Info] Total Bins 97920
[LightGBM] [Info] Number of data points in the train set: 2774, number of used features: 384
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.500000 -> initscore=0.000000
[LightGBM] [Info] Start training from score 0.000000
precision      recall    f1-score   support
0            0.757     0.927     0.833      521
1            0.321     0.104     0.157      173
accuracy                           0.722      694
macro avg       0.539     0.516     0.495      694
weighted avg    0.648     0.722     0.665      694
```

c:\Users\brahi\AppData\Local\Programs\Python\Python313\Lib\site-packages\sklearn\utils\validation.py:2749: UserWarning: X does not have valid feature names, but LGBMClassifier was fitted with feature names
warnings.warn(




(h) Modèle GradientBoost

```
In [102]: from sklearn.ensemble import GradientBoostingClassifier

gb = GradientBoostingClassifier(n_estimators=200, learning_rate=0.1, max_depth=3, random_state=42)
f1_gb = train_evaluate(gb, X_train_emb, y_train, X_test_emb, y_test, "Gradient Boosting")
```

===== Gradient Boosting =====

	precision	recall	f1-score	support
0	0.749	0.944	0.835	521
1	0.216	0.046	0.076	173
accuracy			0.720	694
macro avg	0.483	0.495	0.456	694
weighted avg	0.616	0.720	0.646	694



5-Comparaison et sélection du meilleur modèle

- Modèles utilisés et F1-score associés

```
In [103]: model_preds= {"Logistic Regression": f1_logreg,
                      "SVM": f1_svm,
                      "Random Forest": f1_rf,
                      "Naive Bayes": f1_gnb,
                      "KNN": f1_knn,
                      "XGBoost": f1_xgb,
                      "LightGBM": f1_lgb,
                      "Gradient Boosting": f1_gb
                    }
models={"Logistic Regression": logreg,
        "SVM": svm,
        "Random Forest": rf,
        "Naive Bayes": gnb,
        "KNN": knn,
        "XGBoost": xgb,
        "LightGBM": lgb_model,
        "Gradient Boosting": gb
      }
```

1. Définition de la classe minoritaire

```
In [104]: from collections import Counter

counter = Counter(y_train)
minority_class = min(counter, key=counter.get)
print("Classe minoritaire =", minority_class)
```

Classe minoritaire = 1

1. Calcul des métriques pour la classe minoritaire

```
In [105]: results = []

for name in model_preds.keys():
    y_pred = model_preds[name]
    model = models[name]

    # F1 score par classe (array)
    f1_per_class = f1_score(y_test, y_pred, average=None)
    f1_minority = f1_per_class[minority_class]

    # Precision & Recall pour la classe minoritaire
    precision_per_class = precision_score(y_test, y_pred, average=None)
    recall_per_class = recall_score(y_test, y_pred, average=None)

    precision_minority = precision_per_class[minority_class]
    recall_minority = recall_per_class[minority_class]

    # AUC si possible
    if hasattr(model, "predict_proba"):
        y_proba = model.predict_proba(X_test_emb)[:, 1]
        auc = roc_auc_score(y_test, y_proba)
    elif hasattr(model, "decision_function"):
        y_score = model.decision_function(X_test_emb)
        auc = roc_auc_score(y_test, y_score)
    else:
        auc = None

    results.append({
        "Model": name,
        "Precision_minority": precision_minority,
        "Recall_minority": recall_minority,
        "F1_minority": f1_minority,
        "AUC": auc
    })
```

```
c:\Users\brahi\AppData\Local\Programs\Python\Python313\Lib\site-packages\sklearn\utils\validation.py:2749: UserWarning: X does not have valid feature names
s, but LGBMClassifier was fitted with feature names
warnings.warn(
```

1. Présenter les résultats dans un tableau comparatif

```
In [106]: df_results = pd.DataFrame(results)
df_results = df_results.sort_values(by="F1_minority", ascending=False)
df_results.reset_index(drop=True, inplace=True)
df_results
```

Out[106]:

	Model	Precision_minority	Recall_minority	F1_minority	AUC
0	Logistic Regression	0.299296	0.491329	0.371991	0.567628
1	SVM	0.296820	0.485549	0.368421	0.554447
2	Naive Bayes	0.316062	0.352601	0.333333	0.560061
3	XGBoost	0.307692	0.115607	0.168067	0.514517
4	LightGBM	0.321429	0.104046	0.157205	0.516892
5	KNN	0.288136	0.098266	0.146552	0.538382
6	Gradient Boosting	0.216216	0.046243	0.076190	0.529362
7	Random Forest	0.000000	0.000000	0.000000	0.563922

1. Sélection du meilleur modèle selon le F1 de la classe minoritaire

```
In [107]: best_model = df_results.loc[0, "Model"]
print(f"Best model based on F1 score for minority class: {best_model}")
```

Best model based on F1 score for minority class: Logistic Regression

6-Test final et sauvegarde

1. Chargement et aperçu du fichier de test

```
In [108]: test_data = pd.read_csv('task_A_En_test.csv')
test_data.head()
```

Out[108]:

	tweet	sarcastic
0	Size on the the Toulouse team, That pack is mo...	0
1	Pinball!	0
2	So the Scottish Government want people to get ...	1
3	villainous pro tip : change the device name on...	0
4	I would date any of these men 😊	0

1. Prétraitement des textes du fichier de test

```
In [109]: test_data['text_cleaned'] = test_data['tweet'].apply(clean_text)
test_data.head()
```

Out[109]:

	tweet	sarcastic	text_cleaned
0	Size on the the Toulouse team, That pack is mo...	0	size toulouse team pack monstrous cant see wel...
1	Pinball!	0	pinball
2	So the Scottish Government want people to get ...	1	scottish government want people booster shots ...
3	villainous pro tip : change the device name on...	0	villainous pro tip change device name bluetooth...
4	I would date any of these men 😊	0	date men

1. Encodage des textes du test avec SBERT

```
In [110]: X_test_final_list = test_data['text_cleaned'].astype(str).tolist()
X_test_final_emb = model_sbert.encode(X_test_final_list, batch_size=32, show_p
rogress_bar=True)
```

Batches: 100%|██████████| 44/44 [00:02<00:00, 17.36it/s]

1. Prédiction sur le fichier de test avec le meilleur modèle

```
In [111]: best_model_instance = models[best_model]
y_test_final_pred = best_model_instance.predict(X_test_final_emb)
test_data['sarcastic_pred'] = y_test_final_pred
test_data.head()
```

Out[111]:

	tweet	sarcastic	text_cleaned	sarcastic_pred
0	Size on the the Toulouse team, That pack is mo...	0	size toulouse team pack monstrous cant see wel...	1
1	Pinball!	0	pinball	0
2	So the Scottish Government want people to get ...	1	scottish government want people booster shots ...	0
3	villainous pro tip : change the device name on...	0	villainous pro tip change device name bluetooth...	1
4	I would date any of these men 😊	0	date men	1

1. Évaluation des prédictions sur le fichier de test


```
In [112]: if 'sarcastic' in test_data.columns:  
    y_true = test_data['sarcastic']  
    y_pred = test_data['sarcastic_pred']  
    print("\nClassification Report on Test Data:")  
    print(classification_report(y_true, y_pred, digits=3))  
    try:  
        y_prob = best_model_instance.predict_proba(X_test_final_emb)[:, 1]  
        auc = roc_auc_score(y_true, y_prob)  
        print(f"AUC : {auc:.3f}")  
    except:  
        print("AUC non disponible (le modèle ne supporte pas predict_proba).")  
    cm = confusion_matrix(y_true, y_pred)  
    sns.heatmap(cm, annot=True, fmt="d", cmap="Blues")  
    plt.title("Matrice de confusion - Test final")  
    plt.xlabel("Prédictions")  
    plt.ylabel("Vérités terrain")  
    plt.show()  
else:  
    print("⚠ La colonne 'sarcastic' n'existe pas dans le fichier test.")
```

Classification Report on Test Data:

	precision	recall	f1-score	support
0	0.909	0.565	0.697	1200
1	0.202	0.660	0.309	200
accuracy			0.579	1400
macro avg	0.555	0.613	0.503	1400
weighted avg	0.808	0.579	0.641	1400

AUC : 0.648

Matrice de confusion - Test final



1. Sauvegarde du modèle et des prédictions finales avec les labels

```
In [113]: joblib.dump(best_model_instance, "best_model_sbert.pkl")
test_data[["text_cleaned", "sarcastic_pred", "sarcastic"]].to_csv("final_predictions.csv", index=False)
print("✓ Modèle sauvegardé : best_model_sbert.pkl")
print("✓ Prédictions sauvegardées : final_predictions.csv")
```

✓ Modèle sauvegardé : best_model_sbert.pkl
✓ Prédictions sauvegardées : final_predictions.csv